


Simple Linear Regression Model

```
In [1]: from Utils import *
```

```
In [2]: # Load training data
trainImages, trainOutputxi, trainOutputyi, trainOutputxw, trainOutputyw = readIm
```

```
In [3]: # normalizing pixel values
trainImagesNormalized = np.array(trainImages) / 255.0
```

```
In [4]: # show some training images
titles = [f"x={x}, y={y}" for x, y in zip(trainOutputxi[:16], trainOutputyi[:16])
display_images(trainImages[:16], rows=4, cols=4, titles=titles)
```



```
In [5]: # Flatten the images for the linear regression model
trainImages_flattened = \
[img.flatten() for img in trainImagesNormalized]
trainImages_flattened = np.array(trainImages_flattened)
```

```
In [6]: # Prepare the Outputs
trainOutputs = np.array([
    trainOutputxi, # x values
    trainOutputyi, # y values
    trainOutputxw, # x_width values
    trainOutputyw, # y_width values
], dtype=float).T # Transpose to align correctly
```

```
In [7]: # Split the Data
from sklearn.model_selection import train_test_split

X_train, X_val, y_train, y_val = train_test_split(
    trainImages_flattened, trainOutputs, test_size=0.2, random_state=42
)
```

```
In [8]: # Train the Linear Regression Model
from sklearn.linear_model import LinearRegression

# Initialize the Linear Regression model
lr_model = LinearRegression()

# Train the model
lr_model.fit(X_train, y_train)
```

```
Out[8]: ▼ LinearRegression ⓘ ⓘ
LinearRegression()
```

```
In [9]: # Evaluate the Model
from sklearn.metrics import mean_squared_error

# Predict on the validation set
y_pred = lr_model.predict(X_val)

# Calculate MSE for validation
mse = mean_squared_error(y_val, y_pred)
print(f"Mean Squared Error on Validation Set: {mse:.2f}")
```

Mean Squared Error on Validation Set: 1030.05

```
In [10]: # Save the Model
import joblib

# Save the trained model to a file
joblib.dump(lr_model, "linear_regression_models/linear_regression_simple_model.p
```

```
Out[10]: ['linear_regression_models/linear_regression_simple_model.pkl']
```

Improving the model

```
In [11]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Ridge
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error
```

```
In [12]: # Standardize the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_val_scaled = scaler.transform(X_val)

# Save the scaler
joblib.dump(scaler, "linear_regression_scalers/scaler.pkl")
```

```
Out[12]: ['linear_regression_scalers/scaler.pkl']
```

```
In [13]: # Define the Ridge regression model
ridge_model = Ridge()

# Define the hyperparameter grid for alpha
param_grid = {'alpha': [0.01, 0.1, 1, 10, 100, 1000]} # A range of values for alpha
```

```
In [14]: # Set up GridSearchCV to find the best alpha
grid_search = GridSearchCV(ridge_model, param_grid, cv=5, scoring='neg_mean_squared_error')
```

```
In [15]: # Fit the grid search on the training data
grid_search.fit(X_train_scaled, y_train)
```

Fitting 5 folds for each of 6 candidates, totalling 30 fits

```
Out[15]: GridSearchCV
  ▸ best_estimator_: Ridge
    ▸ Ridge
```

```
In [16]: import joblib

# Save models for each combination of hyperparameters
for idx, alpha_value in enumerate(grid_search.cv_results_['param_alpha']):
    model = Ridge(alpha=alpha_value)
    model.fit(X_train_scaled, y_train)

    model_filename = f"linear_regression_models/ridge_model_alpha_{alpha_value}.pkl"
    joblib.dump(model, model_filename)
    print(f"Model with alpha {alpha_value} saved as {model_filename}")
```

```
Model with alpha 0.01 saved as linear_regression_models/ridge_model_alpha_0.01.pkl
Model with alpha 0.1 saved as linear_regression_models/ridge_model_alpha_0.1.pkl
Model with alpha 1.0 saved as linear_regression_models/ridge_model_alpha_1.0.pkl
Model with alpha 10.0 saved as linear_regression_models/ridge_model_alpha_10.0.pkl
Model with alpha 100.0 saved as linear_regression_models/ridge_model_alpha_100.0.pkl
Model with alpha 1000.0 saved as linear_regression_models/ridge_model_alpha_1000.0.pkl
```

```
In [17]: # Get the best alpha value
best_alpha = grid_search.best_params_['alpha']
print(f"Best alpha found: {best_alpha}")
```

```
# Get the best model
best_model = grid_search.best_estimator_
```

Best alpha found: 1000

```
In [18]: # Predict on the validation set
y_pred = best_model.predict(X_val_scaled)

# Calculate MSE for the validation set
mse = mean_squared_error(y_val, y_pred)
print(f"Mean Squared Error on Validation Set with Best Alpha: {mse:.2f}")
```

Mean Squared Error on Validation Set with Best Alpha: 1027.29

In []: